

6. - Deep Learning APIs

6.2 - Deep Learning using TensorFlow

Rodolfo Sabino

10 Setembro, 2020

Departamento de de Computação



Disclaimer

The following content is heavily based on HCIA-AI Course material by Huawei Technologies Co., Ltd., authored by Zhang Luxiao and Yuan Meng. Distribution is not allowed.



Index

- TensorFlow 2.x Environment Setup
- TensorFlow 2.x Development Overview
- Study Case: MNIST Digit Recognition



TensorFlow 2.x Environment Setup



TensorFlow 2.x Installation

First, install Python3 and pip, maybe using a manager such as [Miniconda](#)

By using **pip** package manager, the installation is simple from command line:

```
pip install tensorflow
```

For using the stable release of TensorFlow for CPU (recommended), or

```
pip install tensorflow-gpu
```

for GPU-acceleration support (avoid this unless you know what you are doing).

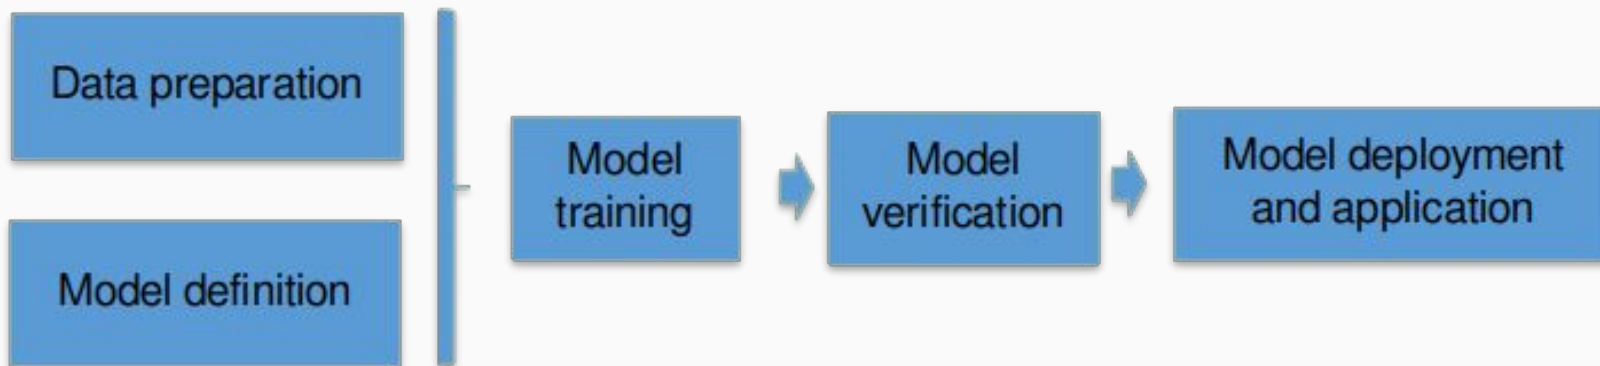
Or simply create a Colaboratory Notebook in Google Drive



TensorFlow 2.x Development Overview



TensorFlow Development Process




Defining the Model

Models in TensorFlow can be either **sequential** or **functional**.


Sequential models provided by [tf.keras.Sequential](https://keras.io/api/sequential/) are simple to use and cover a wide range of common scenarios.

Functional models provided by [tf.keras.Model](https://keras.io/api/models/model_functional/) are more complex allow to create more versatile applications, with non-sequential connections and multiple input / outputs.

1. Sequential API



2. Functional API



[Image source](#)



Defining the Model

Building and configuring layers of a neural network. [tf.keras.layers](https://keras.io/layers/) offers different types of layers based on different algorithms and optimized for different applications, such as image convolution layers and RRN layers.



Training the Model

After defining the model, use `tf.keras.Model.compile()` to compile it, i.e., to configure how the underlying learning process is carried out by defining a loss function, an optimizer, evaluation metrics, and other setups. Common options can be found in [tf.keras.losses](#), [tf.keras.optimizers](#), and [tf.keras.metrics](#) modules.

Then `tf.keras.Model.fit()` to **train the model**.

Training the model may take a long time for complex models and big datasets. Callback functions from [tf.keras.callback](#) should be employed to add breakpoints and give opportunity of [saving and further loading](#) the model training set.



Saving and Loading the Model

The model can be saved during the training as a checkpoint.

A model can be saved directly by invoking [`tf.keras.Model.save\(\)`](#) / [`tf.keras.Model.save_weights\(\)`](#) once the training is over.

Finally, a model can be loaded by calling [`tf.keras.models.load_model\(\)`](#) or [`tf.keras.Model.load_weights\(\)`](#), for example.



Testing the Model

After training, [tf.keras.Model.evaluate\(\)](#) can be used to **test the model**. It returns useful metrics such as loss and accuracy values.



Using the model

Finally, **using the model for inference** is relatively easy. Use [`tf.keras.Model.predict\(\)`](#) with an input tensor acting as samples or batches of samples and this function will generate outputs predictions for the input samples.

Further reading on these steps:


https://www.tensorflow.org/guide/keras/train_and_evaluate



Study Case: MNIST Digit Recognition



MNIST Dataset



Can be downloaded from <http://yann.lecun.com/exdb/mnist/>

Different from printed fonts, handwriting of different people has different sizes and styles, making it difficult for computers to recognize handwriting.

The MNIST datasets consist of a training set and a test set

- Training set: 60,000 handwriting images and corresponding labels
- Test set: 10,000 handwriting images and corresponding labels



MNIST Dataset

Examples



Corresponding labels

[0,0,0,0,0,
1,0,0,0,0]

[0,0,0,0,0,
0,0,0,0,1]

[0,0,0,0,0,
0,0,1,0,0]

[0,0,0,1,0,
0,0,0,0,0]

[0,0,0,0,1,
0,0,0,0,0]



TensorFlow 2.x and MNIST: Quick and Easy

```
import tensorflow as tf
import tensorflow.keras as keras

(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(10)
])

model.compile(optimizer='adam',
              loss=keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

model.fit(x_train, y_train, batch_size=32, epochs=5)
model.evaluate(x_test, y_test, verbose=1)
```



Thank You!

Next: Lab Guides 04 and 05

